

Panic/Anxiety RAG Chatbot - Setup Guide (with RAG-Anything)

Overview

This setup uses **RAG-Anything** for preprocessing + **custom chatbot** for conversation:

Your PDFs → RAG-Anything (ingest.py) → FAISS Vector Store



Custom Chatbot

(safety layer + prompt)

STEP-BY-STEP SETUP

PART A: RAG-Anything Setup (Preprocessing)

Step A1: Clone RAG-Anything Repo

```
bash

# In a working directory
git clone https://github.com/HKUDS/RAG-Anything.git
cd RAG-Anything

# Install dependencies
pip install -r requirements.txt
```

 **Expected time:** 2-3 minutes

Step A2: Prepare Your PDFs

```
bash
```

```
# Create docs directory
```

```
mkdir -p data/docs
```

```
# Copy your 5 trusted mental health PDFs here:
```

```
# - panic_attack_coping.pdf
```

```
# - grounding_exercises.pdf
```

```
# - deep_breathing_manual.pdf
```

```
# - CBT_panic_guide.pdf
```

```
# - WHO_anxiety_tips.pdf
```

```
ls data/docs/ # Verify all 5 PDFs are there
```

Step A3: Run RAG-Anything Preprocessing

```
bash
```

```
# This preprocesses PDFs, creates embeddings, and builds FAISS vector store
```

```
python ingest.py \
```

```
--data_dir ./data/docs \
```

```
--vectorstore faiss \
```

```
--chunk_size 500 \
```

```
--chunk_overlap 50 \
```

```
--embedding_model sentence-transformers/all-MiniLM-L6-v2
```

What happens:

1. ☒ Loads all 5 PDFs from `data/docs/`
2. ☒ Splits into 500-character chunks with 50-char overlap
3. ☒ Creates embeddings using `all-MiniLM-L6-v2`
4. ☒ Builds FAISS index for fast similarity search
5. ☒ Saves to `./vectorstore/` directory

Output files:

```
RAG-Anything/vectorstore/
```

```
├── index.faiss    # Vector index
```

```
└── index.pkl     # Metadata
```

 **Expected time:** 3-5 minutes (first run downloads embedding model)

☒ **You'll see:**

Loading documents...

✓ Loaded 5 PDFs

✓ Created 150 chunks (approx)

✓ Generated embeddings...

✓ Built FAISS index

✓ Saved to ./vectorstore/

PART B: Your Custom Chatbot Setup

Step B1: Create Project Structure

```
bash
```

```
# Go back to parent directory
```

```
cd ..
```

```
# Create anxiety chatbot project
```

```
mkdir anxiety-rag-chatbot
```

```
cd anxiety-rag-chatbot
```

```
# Create directories
```

```
mkdir logs
```

Step B2: Copy Vector Store

```
bash
```

```
# Copy the preprocessed vector store from RAG-Anything
```

```
cp -r ../RAG-Anything/vectorstore ./
```

```
# Verify it's there
```

```
ls vectorstore/
```

```
# Should show: index.faiss index.pkl
```

Step B3: Create Virtual Environment

```
bash
```

```
# Create venv
python -m venv venv

# Activate it
# On Linux/Mac:
source venv/bin/activate

# On Windows:
venv\Scripts\activate
```

Step B4: Install Dependencies

Create `requirements.txt`:

```
langchain==0.1.0
sentence-transformers==2.2.2
faiss-cpu==1.7.4
torch==2.0.0
transformers==4.35.0
huggingface-hub==0.19.0
```

Then install:

```
bash

pip install -r requirements.txt
```

 **Expected time:** 5-10 minutes

Step B5: Place Chatbot Script

Save the main chatbot code as `chatbot.py` in your project root

Step B6: Verify Setup

```
bash

# Check that everything is in place
ls -la

# Should show:
# - chatbot.py
# - requirements.txt
# - vectorstore/
# - logs/
# - venv/
```



RUNNING THE CHATBOT






First Launch

```
bash

# Make sure venv is activated
source venv/bin/activate # or on Windows: venv\Scripts\activate

# Run chatbot
python chatbot.py
```

What happens:

1.  Validates vector store exists
2.  Loads embeddings model (cached)
3.  Initializes FAISS vector store
4.  Loads LLM
5.  Starts interactive chat



Startup time: 30-60 seconds

You'll see:

PANIC & ANXIETY RAG CHATBOT - STARTUP

✓ Logs directory ready

✓ Vector store validated at ./vectorstore

📁 Loading vector store...

✓ Vector store loaded successfully

Embedding model: sentence-transformers/all-MiniLM-L6-v2

Total vectors indexed: 150 (approx)

🤖 Initializing LLM...

Using HuggingFace (gpt2)

✓ LLM initialized

🎯 Creating RAG chain...

✓ RAG chain created

💡 PANIC & ANXIETY SUPPORT CHATBOT

Powered by Trusted Mental Health Resources

You: _

Example Interaction

You: I'm having a panic attack right now, help me

🕒 Thinking...

🤖 Assistant:

I'm glad you reached out. Let's work through this together using a proven technique.

Here's what we'll do - Box Breathing:

1. Breathe in slowly for 4 counts
2. Hold your breath for 4 counts
3. Breathe out slowly for 4 counts
4. Hold for 4 counts

Repeat 5-10 times

This activates your parasympathetic nervous system, which calms your body...

📄 Sources:

- deep_breathing_manual.pdf
- panic_attack_coping.pdf

💡 Remember: For persistent symptoms, please consult a mental health professional.

You: _

⚙️ CONFIGURATION

Switch to Ollama (Recommended for Better Quality)

Install Ollama:

```
bash

# Visit: https://ollama.ai
# Download and install for your OS
```

In your terminal (separate window):

```
bash

# Pull and run Mistral (best quality)
ollama pull mistral
ollama serve
```

In `chatbot.py`, change:

```
python

class Config:
    USE_OLLAMA = True # Change from False to True
    OLLAMA_MODEL = "mistral" # or "neural-chat", "orca-mini"
```

Restart chatbot:

```
bash

python chatbot.py
```

Quality comparison:

Model	Speed	Quality	Use Case
GPT-2	⚡ Fast (2-3s)	Basic	Testing, CPU-limited
Mistral (Ollama)	🕒 Medium (5-10s)	Excellent	Production, recommended
Neural-chat (Ollama)	🕒 Fast (3-5s)	Very Good	Balance of speed & quality

"Vector store not found"

✗ Vector store not found at ./vectorstore

Fix:

```
bash

# Make sure you copied vectorstore from RAG-Anything
cp -r ../RAG-Anything/vectorstore ./

# Verify
ls vectorstore/index.faiss #Should exist
```

"No module named 'langchain'"

```
bash

pip install --upgrade pip
pip install -r requirements.txt
```

"Slow startup"

- **First run:** Normal (downloads models)
- **Subsequent runs:** Should be 30-60 seconds
- **If still slow:** Close other apps, check disk space

"Out of memory"

```
python

# In chatbot.py, reduce TOP_K:
TOP_K = 2 # Changed from 3
```

Or use smaller embedding model:

```
python

EMBEDDING_MODEL = "sentence-transformers/all-MiniLM-L6-v2" #Already minimal
```


Ollama connection error

Error connecting to Ollama

Fix:

```
bash
```

```
# In separate terminal, make sure Ollama is running:
```

```
ollama serve
```

```
# Then run chatbot in another terminal
```

```
python chatbot.py
```



PERFORMANCE METRICS

First Run (RAG-Anything preprocessing)

- **Time:** 3-5 minutes
- **What:** PDF loading, chunking, embedding, FAISS indexing
- **Output:** 150+ chunks indexed

Chatbot Startup

- **Time:** 30-60 seconds
- **What:** Load embeddings, vector store, LLM
- **One-time:** These are cached after first load

Per Query

- **Time:** 2-10 seconds (depends on LLM)
 - FAISS search: ~0.1s
 - LLM generation: 2-9s (GPT-2) or 5-10s (Mistral)
- **Memory:** ~1-2GB



FINAL PROJECT STRUCTURE

```
anxiety-rag-chatbot/
```

```
|— venv/                # Virtual environment
```

```
|— vectorstore/         # Pre-built FAISS index (from RAG-Anything)
```

	index.faiss	# Vector embeddings
	index.pkl	# Metadata
	logs/	# Conversation logs (optional)
	chatbot.py	# Main chatbot script ← RUN THIS
	requirements.txt	# Python dependencies
	README.md	# Documentation (optional)

QUICK START CHECKLIST

- ☐ Cloned RAG-Anything
- ☐ Placed 5 PDFs in `RAG-Anything/data/docs/`
- ☐ Ran `python ingest.py` successfully
- ☐ Copied `vectorstore/` to `anxiety-rag-chatbot/`
- ☐ Created Python venv
- ☐ Installed dependencies from `requirements.txt`
- ☐ Placed `chatbot.py` in project root
- ☐ Verified `vectorstore/index.faiss` exists
- ☐ First run: `python chatbot.py`
- ☐ Asked test question: "How do I breathe calmly?"
- ☐ Got response with source documents

YOU'RE READY!

```
bash

# Activate venv
source venv/bin/activate

# Run chatbot
python chatbot.py

# Start chatting!
You: I feel anxious, help me
```

CRISIS RESOURCES (Built-in)

If you mention self-harm, suicide, or crisis keywords, the chatbot automatically provides:

- **IN India:** AASRA (1800-599-0019), iCall (9152987821)
- **us US:** Crisis Text Line (text HELLO to 741741)

-  **Other:** Localized helplines
-

OPTIONAL ENHANCEMENTS

Add Streamlit Web UI

```
bash

pip install streamlit
# Create app.py with Streamlit wrapper
streamlit run app.py
```

Enable Voice

```
bash

pip install pytsx3 SpeechRecognition
# Add voice input/output
```

Conversation History

```
python

# Save/load conversations as JSON
# Add to run_chatbot() function
```

Questions? Issues? Check troubleshooting section above or run with `--debug` flag 